

Asg 19.6

Working with Pandas DataFrames -- The Essentials (Part Four)

(Coding)



Files needed for this assignment:

marketing_campaign.csv
baseball_data.csv
student.txt
cereals.txt

```
In [1]: # set up notebook to display multiple output in one cell

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

print('The notebook is set up to display multiple output in one cell.')
```

The notebook is set up to display multiple output in one cell.

```
In [2]: # conventional way to import pandas and numpy

import pandas as pd
import numpy as np
```

PART ONE

<div class="alert alert-block alert-info"

For Questions 1-9: We will be using the '**marketing_campaign.csv**' dataset and the **customers DataFrame** </div>

Question 1:

a. Read in the dataset '**marketing_campaign.csv**' and store the results in a DataFrame named **customers**.

b. Use appropriate attributes and methods to inspect the **customers** DataFrame.

```
In [4]: customers = pd.read_csv('marketing_campaign.csv', sep=';')
print(customers.head())
print(customers.tail())
print(customers.info())
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	\
0	5524	1957	Graduation	Single	58138.0	0	0	
1	2174	1954	Graduation	Single	46344.0	1	1	
2	4141	1965	Graduation	Together	71613.0	0	0	
3	6182	1984	Graduation	Together	26646.0	1	0	
4	5324	1981	PhD	Married	58293.0	1	0	

	Dt_Customer	Recency	MntWines	...	NumWebVisitsMonth	AcceptedCmp3	\
0	2012-09-04	58	635	...	7	0	
1	2014-03-08	38	11	...	5	0	
2	2013-08-21	26	426	...	4	0	
3	2014-02-10	26	11	...	6	0	
4	2014-01-19	94	173	...	5	0	

	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Z_CostContact	Z_Revenue	Response
0	3	11	1
1	3	11	0
2	3	11	0
3	3	11	0
4	3	11	0

[5 rows x 29 columns]

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	\
2235	10870	1967	Graduation	Married	61223.0	0	
2236	4001	1946	PhD	Together	64014.0	2	
2237	7270	1981	Graduation	Divorced	56981.0	0	
2238	8235	1956	Master	Together	69245.0	0	
2239	9405	1954	PhD	Married	52869.0	1	

	Teenhome	Dt_Customer	Recency	MntWines	...	NumWebVisitsMonth	\
2235	1	2013-06-13	46	709	...	5	
2236	1	2014-06-10	56	406	...	7	
2237	0	2014-01-25	91	908	...	6	
2238	1	2014-01-24	8	428	...	3	
2239	1	2012-10-15	40	84	...	7	

	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	\
2235	0	0	0	0	0	
2236	0	0	0	1	0	
2237	0	1	0	0	0	
2238	0	0	0	0	0	
2239	0	0	0	0	0	

	Complain	Z_CostContact	Z_Revenue	Response
2235	0	3	11	0
2236	0	3	11	0
2237	0	3	11	0
2238	0	3	11	0
2239	0	3	11	1

[5 rows x 29 columns]

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239

Data columns (total 29 columns):

#	Column	Non-Null Count	Dtype
0	ID	2240 non-null	int64
1	Year_Birth	2240 non-null	int64
2	Education	2240 non-null	object
3	Marital_Status	2240 non-null	object
4	Income	2216 non-null	float64
5	Kidhome	2240 non-null	int64
6	Teenhome	2240 non-null	int64
7	Dt_Customer	2240 non-null	object
8	Recency	2240 non-null	int64
9	MntWines	2240 non-null	int64
10	MntFruits	2240 non-null	int64
11	MntMeatProducts	2240 non-null	int64
12	MntFishProducts	2240 non-null	int64
13	MntSweetProducts	2240 non-null	int64
14	MntGoldProds	2240 non-null	int64
15	NumDealsPurchases	2240 non-null	int64
16	NumWebPurchases	2240 non-null	int64
17	NumCatalogPurchases	2240 non-null	int64
18	NumStorePurchases	2240 non-null	int64
19	NumWebVisitsMonth	2240 non-null	int64
20	AcceptedCmp3	2240 non-null	int64
21	AcceptedCmp4	2240 non-null	int64
22	AcceptedCmp5	2240 non-null	int64
23	AcceptedCmp1	2240 non-null	int64
24	AcceptedCmp2	2240 non-null	int64
25	Complain	2240 non-null	int64
26	Z_CostContact	2240 non-null	int64
27	Z_Revenue	2240 non-null	int64
28	Response	2240 non-null	int64

dtypes: float64(1), int64(25), object(3)

memory usage: 507.6+ KB

None

Question 2:

Filter the **customers** DataFrame rows to show only customers with an '**Income**' of at least 100000.

```
In [6]: customers.loc[customers['Income']>=1000,:]
```

```
Out[6]:
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer	R
0	5524	1957	Graduation	Single	58138.0	0	0	2012-09-04	
1	2174	1954	Graduation	Single	46344.0	1	1	2014-03-08	
2	4141	1965	Graduation	Together	71613.0	0	0	2013-08-21	
3	6182	1984	Graduation	Together	26646.0	1	0	2014-02-10	
4	5324	1981	PhD	Married	58293.0	1	0	2014-01-19	
...
2235	10870	1967	Graduation	Married	61223.0	0	1	2013-06-13	
2236	4001	1946	PhD	Together	64014.0	2	1	2014-06-10	
2237	7270	1981	Graduation	Divorced	56981.0	0	0	2014-01-25	
2238	8235	1956	Master	Together	69245.0	0	1	2014-01-24	
2239	9405	1954	PhD	Married	52869.0	1	1	2012-10-15	

2216 rows × 29 columns

Question 3:

Select the **'Education'** Series from the Question 2 filtered DataFrame using the `.loc` label-based indexing property.

Note: See the image below for information on selecting data via the Pandas `.loc` (label-based indexing) and `.iloc` (position-based indexing) indexing properties: !

[loc%20and%20iloc.jpg](attachment:loc%20and%20iloc.jpg)

Documentation for **'loc'**

Documentation for **'iloc'**

```
In [10]: customers.loc[:, 'Education']
```

```
Out[10]:
```

0	Graduation
1	Graduation
2	Graduation
3	Graduation
4	PhD
...	
2235	Graduation
2236	PhD
2237	Graduation
2238	Master
2239	PhD

Name: Education, Length: 2240, dtype: object

Question 4:

Select the **'Marital_Status'** Series from the Question 2 filtered DataFrame without using the .loc label-based indexing attribute.

```
In [11]: customers['Marital_Status']
```

```
Out[11]: 0      Single
1      Single
2    Together
3    Together
4      Married
...
2235   Married
2236   Together
2237   Divorced
2238   Together
2239   Married
Name: Marital_Status, Length: 2240, dtype: object
```

Question 5:

Filter the **customers** DataFrame rows to show only customers with an **'Income'** of at least 80000 and a **'Marital_Status'** of Single.


```
In [21]: customers.loc[(customers['Income']>=80000) & (customers['Marital_Status']=='Single']
```

Out[21]:

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer
15	2114	1946	PhD	Single	82800.0	0	0	2012-11-24
67	9369	1979	2n Cycle	Single	88194.0	0	1	2014-03-15
113	999	1991	Graduation	Single	86037.0	0	0	2013-01-03
124	7215	1983	Graduation	Single	101970.0	0	0	2013-03-12
140	821	1992	Master	Single	92859.0	0	0	2012-10-15
159	2730	1955	Graduation	Single	80317.0	0	0	2013-08-20
418	4216	1981	Graduation	Single	91065.0	0	0	2013-02-22
430	3725	1961	PhD	Single	84865.0	0	0	2013-05-05
447	1137	1964	Graduation	Single	81246.0	0	0	2013-12-25
456	4947	1966	2n Cycle	Single	89572.0	0	0	2012-09-15
507	6071	1989	Graduation	Single	81217.0	0	0	2013-07-15
515	203	1975	Master	Single	81169.0	0	0	2013-04-14
561	3179	1980	Graduation	Single	81741.0	0	0	2013-06-16
626	10156	1975	Graduation	Single	84196.0	0	1	2013-06-03
634	8923	1973	Graduation	Single	83917.0	0	0	2013-04-18
636	6945	1952	Graduation	Single	84574.0	0	0	2013-06-04
650	4248	1960	Master	Single	98777.0	0	0	2014-02-17
686	9826	1972	PhD	Single	86857.0	0	0	2012-09-12
703	8029	1988	Master	Single	90247.0	0	0	2014-04-25
803	9930	1944	PhD	Single	82716.0	0	0	2013-11-05
878	1446	1956	Master	Single	86424.0	0	0	2014-04-05
884	5830	1972	PhD	Single	86857.0	0	0	2012-09-12
905	11074	1977	Graduation	Single	85072.0	0	0	2014-04-05
906	10150	1961	Graduation	Single	86429.0	0	0	2013-11-27
914	10619	1994	Graduation	Single	95529.0	0	0	2012-12-03
1001	7962	1987	PhD	Single	95169.0	0	0	2013-10-05
1031	9220	1971	Graduation	Single	91700.0	0	0	2013-01-17
1066	3005	1992	Graduation	Single	83528.0	0	0	2014-05-07
1097	10245	1986	2n Cycle	Single	80910.0	0	0	2012-10-31
1111	1524	1983	2n Cycle	Single	81698.0	0	0	2013-03-07
1113	7451	1960	Master	Single	98777.0	0	0	2014-02-17
1126	6749	1966	Graduation	Single	86358.0	1	1	2012-08-08
1169	2410	1969	Graduation	Single	81657.0	0	0	2014-01-22
1179	5735	1991	Master	Single	90638.0	0	0	2014-02-13
1269	9400	1958	2n Cycle	Single	85485.0	0	0	2014-06-27

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer
1333	5147	1948	Graduation	Single	90842.0	0	0	2013-07-25
1358	1065	1963	Graduation	Single	80695.0	0	0	2013-12-25
1458	10133	1970	Graduation	Single	93790.0	0	0	2014-02-15
1520	4278	1983	PhD	Single	87188.0	0	0	2013-06-05
1570	4261	1946	PhD	Single	82800.0	0	0	2012-11-25
1572	5350	1991	Master	Single	90638.0	0	0	2014-02-15
1651	8395	1961	Graduation	Single	82014.0	0	0	2012-08-25
1689	295	1989	Graduation	Single	81217.0	0	0	2013-07-15
1722	569	1991	Graduation	Single	90273.0	0	0	2013-12-15
1850	4427	1995	2n Cycle	Single	83257.0	0	0	2012-09-15
1854	10163	1984	PhD	Single	82733.0	0	0	2013-09-15
1897	5558	1954	PhD	Single	90933.0	0	0	2014-03-31
1898	4619	1945	PhD	Single	113734.0	0	0	2014-05-25
1922	3138	1956	Graduation	Single	91249.0	0	0	2012-10-25
1924	7966	1959	Graduation	Single	80982.0	1	1	2013-01-05
1952	3434	1951	Graduation	Single	80872.0	0	0	2014-05-15
1958	2109	1990	Graduation	Single	96843.0	0	0	2013-04-25
1992	6248	1947	Master	Single	91712.0	0	0	2013-10-15
1993	10164	1958	Graduation	Single	94472.0	0	1	2014-04-05
2041	4974	1970	Graduation	Single	83273.0	1	2	2012-09-25
2109	3104	1961	Graduation	Single	82332.0	0	0	2012-09-15
2167	3520	1990	Master	Single	91172.0	0	0	2013-03-25
2190	4418	1983	Master	Single	89616.0	0	0	2013-02-25
2213	3661	1995	2n Cycle	Single	80617.0	0	0	2012-10-15

59 rows × 9 columns



Question 6:


Filter the **customers** DataFrame rows to show only customers with an '**Income**' greater than 80000, a '**Marital_Status**' of Divorced, and an '**Education**' level of Graduation.

```
In [20]: customers.loc[(customers['Income']>=80000) & (customers['Marital_Status']=='Divo
```


Out[20]:

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Custom
53	2225	1977	Graduation	Divorced	82582.0	0	0	2014-06-
84	535	1987	Graduation	Divorced	81361.0	0	0	2014-02-
252	10089	1974	Graduation	Divorced	102692.0	0	0	2013-04-
345	1411	1952	Graduation	Divorced	82623.0	0	0	2013-11-
347	1826	1970	Graduation	Divorced	84835.0	0	0	2014-06-
655	5555	1975	Graduation	Divorced	153924.0	0	0	2014-02-
734	10430	1973	Graduation	Divorced	89694.0	1	1	2013-10-
837	5687	1980	Graduation	Divorced	81702.0	0	0	2012-09-
942	6810	1983	Graduation	Divorced	82025.0	0	0	2013-05-
990	8545	1954	Graduation	Divorced	85683.0	0	0	2014-03-
1100	5538	1975	Graduation	Divorced	83829.0	0	0	2013-10-
1265	3910	1975	Graduation	Divorced	83829.0	0	0	2013-10-
1582	3503	1950	Graduation	Divorced	82460.0	0	0	2013-12-
1584	4608	1987	Graduation	Divorced	81361.0	0	0	2014-02-

14 rows × 29 columns



Question 7:

Filter the **customers** DataFrame rows to show only customers with a **'Marital_Status'** of Married or Together or an **'Education'** level of PhD.


In [22]: `customers.loc[(customers['Education']=='PhD') | (customers['Marital_Status']=='`

Out[22]:

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer	
	2	4141	1965	Graduation	Together	71613.0	0	0	2013-08
	3	6182	1984	Graduation	Together	26646.0	1	0	2014-08
	4	5324	1981	PhD	Married	58293.0	1	0	2014-08
	5	7446	1967	Master	Together	62513.0	0	1	2013-08
	7	6177	1985	PhD	Married	33454.0	1	0	2013-08

	2234	8372	1974	Graduation	Married	34421.0	1	0	2013-08
	2235	10870	1967	Graduation	Married	61223.0	0	1	2013-08
	2236	4001	1946	PhD	Together	64014.0	2	1	2014-08
	2238	8235	1956	Master	Together	69245.0	0	1	2014-08
	2239	9405	1954	PhD	Married	52869.0	1	1	2012-10

1621 rows × 29 columns



Question 8:

Filter the **customers** DataFrame rows to show only customers with an **Income** between 40000 and 60000 or an **"Education"** level of Master or PhD.


In [24]: `customers.loc[((customers['Income']>=40000) & (customers['Income']<=60000)) |`

Out[24]:

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer	
	0	5524	1957	Graduation	Single	58138.0	0	0	2012-08-01
	1	2174	1954	Graduation	Single	46344.0	1	1	2014-08-01
	4	5324	1981	PhD	Married	58293.0	1	0	2014-08-01
	5	7446	1967	Master	Together	62513.0	0	1	2013-08-01
	6	965	1971	Graduation	Divorced	55635.0	0	1	2012-10-01

	2227	4201	1962	Graduation	Single	57967.0	0	1	2013-08-01
	2231	9817	1970	Master	Single	44802.0	0	0	2012-08-01
	2237	7270	1981	Graduation	Divorced	56981.0	0	0	2014-08-01
	2238	8235	1956	Master	Together	69245.0	0	1	2014-08-01
	2239	9405	1954	PhD	Married	52869.0	1	1	2012-10-01

888 rows × 29 columns



Question 9:

Use the **"isin() method"** to filter the **customers** DataFrame to show only customers with a **Marital_Status** of Single, Married, or Together.

Documentation for **'isin'**

```
In [31]: specific_statuses = customers['Marital_Status'].isin(['Single', 'Married', 'Together'])
customers[specific_statuses]
```

```
Out[31]:
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer
0	5524	1957	Graduation	Single	58138.0	0	0	201
1	2174	1954	Graduation	Single	46344.0	1	1	201
2	4141	1965	Graduation	Together	71613.0	0	0	201
3	6182	1984	Graduation	Together	26646.0	1	0	201
4	5324	1981	PhD	Married	58293.0	1	0	201
...
2234	8372	1974	Graduation	Married	34421.0	1	0	201
2235	10870	1967	Graduation	Married	61223.0	0	1	201
2236	4001	1946	PhD	Together	64014.0	2	1	201
2238	8235	1956	Master	Together	69245.0	0	1	201
2239	9405	1954	PhD	Married	52869.0	1	1	201

1924 rows × 29 columns

PART TWO

<div class="alert alert-block alert-info">

For Questions 10-14: We will be using the **'baseball_data.csv'** dataset and the **baseball DataFrame** </div>

Question 10:

a. Read in the dataset **'baseball_data.csv'** and store the results in a DataFrame named **baseball**.

See the link below for access to the dataset.

baseball_data.csv

b. Use appropriate attributes and methods to inspect the **baseball** DataFrame.

```
In [34]: baseball = pd.read_csv('baseball_data.csv')
print(baseball.head())
print(baseball.tail())
print(baseball.info())
```

	playerID	yearID	stint	teamID	lgID	G	AB	R	H	2B	...	RBI
SB \												
0	abercda01	1871	1	TRO	NaN	1	4	0	0	0	...	0.0
0.0												
1	addybo01	1871	1	RC1	NaN	25	118	30	32	6	...	13.0
8.0												
2	allisar01	1871	1	CL1	NaN	29	137	28	40	4	...	19.0
3.0												
3	allisdo01	1871	1	WS3	NaN	27	133	28	44	10	...	27.0
1.0												
4	ansonca01	1871	1	RC1	NaN	25	120	29	39	11	...	16.0
6.0												

	CS	BB	SO	IBB	HBP	SH	SF	GIDP
0	0.0	0	0.0	NaN	NaN	NaN	NaN	0.0
1	1.0	4	0.0	NaN	NaN	NaN	NaN	0.0
2	1.0	2	5.0	NaN	NaN	NaN	NaN	1.0
3	1.0	0	2.0	NaN	NaN	NaN	NaN	0.0
4	2.0	2	1.0	NaN	NaN	NaN	NaN	0.0

[5 rows x 22 columns]

	playerID	yearID	stint	teamID	lgID	G	AB	R	H	2B	...	RBI
\												
108784	zimmebr02	2020	1	BAL	AL	2	0	0	0	0	...	0.0
108785	zimmejo02	2020	1	DET	AL	3	0	0	0	0	...	0.0
108786	zimmeky01	2020	1	KCA	AL	16	0	0	0	0	...	0.0
108787	zuberty01	2020	1	KCA	AL	23	0	0	0	0	...	0.0
108788	zuninmi01	2020	1	TBA	AL	28	75	8	11	4	...	10.0

	SB	CS	BB	SO	IBB	HBP	SH	SF	GIDP
108784	0.0	0.0	0	0.0	0.0	0.0	0.0	0.0	0.0
108785	0.0	0.0	0	0.0	0.0	0.0	0.0	0.0	0.0
108786	0.0	0.0	0	0.0	0.0	0.0	0.0	0.0	0.0
108787	0.0	0.0	0	0.0	0.0	0.0	0.0	0.0	0.0
108788	0.0	0.0	6	37.0	0.0	3.0	0.0	0.0	0.0

[5 rows x 22 columns]

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 108789 entries, 0 to 108788
Data columns (total 22 columns):
```

#	Column	Non-Null Count	Dtype
0	playerID	108789 non-null	object
1	yearID	108789 non-null	int64
2	stint	108789 non-null	int64
3	teamID	108789 non-null	object
4	lgID	108052 non-null	object
5	G	108789 non-null	int64
6	AB	108789 non-null	int64
7	R	108789 non-null	int64
8	H	108789 non-null	int64
9	2B	108789 non-null	int64
10	3B	108789 non-null	int64
11	HR	108789 non-null	int64
12	RBI	108033 non-null	float64
13	SB	106421 non-null	float64
14	CS	85248 non-null	float64
15	BB	108789 non-null	int64
16	SO	106689 non-null	float64
17	IBB	72139 non-null	float64

```
18 HBP      105973 non-null float64
19 SH       102721 non-null float64
20 SF       72686 non-null float64
21 GDP      83348 non-null float64
dtypes: float64(9), int64(10), object(3)
memory usage: 18.3+ MB
None
```

Question 11:

a. Read in the dataset **'baseball_data.csv'** again, but this time specify by **name** that the only columns you want to include are the 'R', 'H', 'HR', and 'RBI' columns. Store the results in a DataFrame named **baseball_2**.

b. Use appropriate attributes or methods to check that you read in the correct columns.

```
In [37]: baseball_2 = pd.read_csv('baseball_data.csv', usecols=['R', 'H', 'HR', 'RBI'])
baseball_2
```

```
Out[37]:
```

	R	H	HR	RBI
0	0	0	0	0.0
1	30	32	0	13.0
2	28	40	0	19.0
3	28	44	2	27.0
4	29	39	0	16.0
...
108784	0	0	0	0.0
108785	0	0	0	0.0
108786	0	0	0	0.0
108787	0	0	0	0.0
108788	8	11	4	10.0

108789 rows × 4 columns

Question 12:

a. Read in the dataset **'baseball_data.csv'** again, but this time specify by **position** that the only columns you want to include are the 'R', 'H', 'HR', and 'RBI' columns. Store the results in a DataFrame named **baseball_3**.

b. Use appropriate attributes or methods to check that you read in the correct columns.

```
In [42]: baseball_3 = pd.read_csv('baseball_data.csv',usecols=[7,8,11,12])
baseball_3
```

```
Out[42]:
```

	R	H	HR	RBI
0	0	0	0	0.0
1	30	32	0	13.0
2	28	40	0	19.0
3	28	44	2	27.0
4	29	39	0	16.0
...
108784	0	0	0	0.0
108785	0	0	0	0.0
108786	0	0	0	0.0
108787	0	0	0	0.0
108788	8	11	4	10.0

108789 rows × 4 columns

Question 13:

a. Read in the dataset '**baseball_data.csv**' again, but this time specify that you only want to read in the first 7 rows. Store the results in a DataFrame named **baseball_4**.

b. Use appropriate attributes or methods to check that you read in the correct rows.

```
In [44]: baseball_4 = pd.read_csv('baseball_data.csv',nrows=7)
baseball_4
```

```
Out[44]:
```

	playerID	yearID	stint	teamID	lgID	G	AB	R	H	2B	...	RBI	SB	CS
0	abercda01	1871	1	TRO	NaN	1	4	0	0	0	...	0	0	0
1	addybo01	1871	1	RC1	NaN	25	118	30	32	6	...	13	8	1
2	allisar01	1871	1	CL1	NaN	29	137	28	40	4	...	19	3	1
3	allisdo01	1871	1	WS3	NaN	27	133	28	44	10	...	27	1	1
4	ansonca01	1871	1	RC1	NaN	25	120	29	39	11	...	16	6	2
5	armstbo01	1871	1	FW1	NaN	12	49	9	11	2	...	5	0	1
6	barkeal01	1871	1	RC1	NaN	1	4	0	1	0	...	2	0	0

7 rows × 22 columns



Question 13:

Print out the **teamID** column by iterating through the **'teamID'** Series of the **baseball_4** DataFrame.

```
In [45]: baseball_4['teamID']
```

```
Out[45]:
```

0	TRO
1	RC1
2	CL1
3	WS3
4	RC1
5	FW1
6	RC1

Name: teamID, dtype: object

Question 14:

Use the **iterrows()** method to print out each row of the **baseball_4** DataFrame as a Series.

```
In [52]: gen = (baseball_4.iterrows())
for i in range(len(baseball_4)):
    print(next(gen)[1])
```


playerID	abercda01
yearID	1871
stint	1
teamID	TRO
lgID	NaN
G	1
AB	4
R	0
H	0
2B	0
3B	0
HR	0
RBI	0
SB	0
CS	0
BB	0
SO	0
IBB	NaN
HBP	NaN
SH	NaN
SF	NaN
GIDP	0

Name: 0, dtype: object

playerID	addybo01
yearID	1871
stint	1
teamID	RC1
lgID	NaN
G	25
AB	118
R	30
H	32
2B	6
3B	0
HR	0
RBI	13
SB	8
CS	1
BB	4
SO	0
IBB	NaN
HBP	NaN
SH	NaN
SF	NaN
GIDP	0

Name: 1, dtype: object

playerID	allisar01
yearID	1871
stint	1
teamID	CL1
lgID	NaN
G	29
AB	137
R	28
H	40
2B	4
3B	5
HR	0
RBI	19
SB	3

CS	1
BB	2
SO	5
IBB	NaN
HBP	NaN
SH	NaN
SF	NaN
GIDP	1

Name: 2, dtype: object

playerID	allisdo01
yearID	1871
stint	1
teamID	WS3
lgID	NaN
G	27
AB	133
R	28
H	44
2B	10
3B	2
HR	2
RBI	27
SB	1
CS	1
BB	0
SO	2
IBB	NaN
HBP	NaN
SH	NaN
SF	NaN
GIDP	0

Name: 3, dtype: object

playerID	ansonca01
yearID	1871
stint	1
teamID	RC1
lgID	NaN
G	25
AB	120
R	29
H	39
2B	11
3B	3
HR	0
RBI	16
SB	6
CS	2
BB	2
SO	1
IBB	NaN
HBP	NaN
SH	NaN
SF	NaN
GIDP	0

Name: 4, dtype: object

playerID	armstbo01
yearID	1871
stint	1
teamID	FW1
lgID	NaN

```

G          12
AB         49
R           9
H          11
2B          2
3B          1
HR          0
RBI         5
SB          0
CS          1
BB          0
SO          1
IBB         NaN
HBP         NaN
SH          NaN
SF          NaN
GIDP        0
Name: 5, dtype: object
playerID    barkeal01
yearID      1871
stint        1
teamID      RC1
lgID        NaN
G           1
AB           4
R            0
H            1
2B           0
3B           0
HR           0
RBI          2
SB           0
CS           0
BB           1
SO           0
IBB         NaN
HBP         NaN
SH          NaN
SF          NaN
GIDP        0
Name: 6, dtype: object

```

Documentation for [iterrows](#).

PART THREE

<div class="alert alert-block alert-info"

For Questions 15-16: We will be using the '**student.txt**' **dataset** and the **student DataFrame** </div>

Question 15:

a. Read in the dataset '**student.txt**' and store the results in a DataFrame named **student**.

See the link below for access to the dataset.

[student.txt](#)

b. Use appropriate attributes and methods to inspect the **student** DataFrame.

```
In [55]: student = pd.read_table('student.txt', sep=';')
print(student.head())
print(student.tail())
print(student.info())
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob
Fjob	...	\							
0	GP	F	18	U	GT3	A	4	4	at_home t
teacher	...								
1	GP	F	17	U	GT3	T	1	1	at_home
other	...								
2	GP	F	15	U	LE3	T	1	1	at_home
other	...								
3	GP	F	15	U	GT3	T	4	2	health se
rvices	...								
4	GP	F	16	U	GT3	T	3	3	other
other	...								

	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
0	4	3	4	1	1	3	4	0	11	11
1	5	3	3	1	1	3	2	9	11	11
2	4	3	2	2	3	3	6	12	13	12
3	3	2	2	1	1	5	0	14	14	14
4	4	3	2	1	2	5	0	11	13	13

[5 rows x 33 columns]

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob
Fjob	\								
644	MS	F	19	R	GT3	T	2	3	services
other									
645	MS	F	18	U	LE3	T	3	1	teacher
services									
646	MS	F	18	U	GT3	T	1	1	other
other									
647	MS	M	17	U	LE3	T	3	1	services
services									
648	MS	M	18	R	LE3	T	3	2	services
other									

	...	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
2	G3										
644	...	5	4	2	1	2	5	4	10	1	
1	10										
645	...	4	3	4	1	1	1	4	15	1	
5	16										
646	...	1	1	1	1	1	5	6	11	1	
2	9										
647	...	2	4	5	3	4	2	6	10	1	
0	10										
648	...	4	4	1	3	4	5	4	10	1	
1	11										

[5 rows x 33 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 649 entries, 0 to 648

Data columns (total 33 columns):

#	Column	Non-Null Count	Dtype
0	school	649 non-null	object
1	sex	649 non-null	object
2	age	649 non-null	int64
3	address	649 non-null	object
4	famsize	649 non-null	object
5	Pstatus	649 non-null	object
6	Medu	649 non-null	int64

```

7  Fedu          649 non-null  int64
8  Mjob          649 non-null  object
9  Fjob          649 non-null  object
10 reason        649 non-null  object
11 guardian      649 non-null  object
12 traveltime    649 non-null  int64
13 studytime     649 non-null  int64
14 failures      649 non-null  int64
15 schoolsup      649 non-null  object
16 famsup        649 non-null  object
17 paid          649 non-null  object
18 activities     649 non-null  object
19 nursery       649 non-null  object
20 higher        649 non-null  object
21 internet      649 non-null  object
22 romantic      649 non-null  object
23 famrel        649 non-null  int64
24 freetime      649 non-null  int64
25 goout         649 non-null  int64
26 Dalc          649 non-null  int64
27 Walc          649 non-null  int64
28 health        649 non-null  int64
29 absences      649 non-null  int64
30 G1            649 non-null  int64
31 G2            649 non-null  int64
32 G3            649 non-null  int64
dtypes: int64(16), object(17)
memory usage: 167.4+ KB
None

```

Question 16:

a. Read in the dataset **'student.txt'** again, but use the **select_dtypes() method** to include only the numeric columns. Store the results in a DataFrame named **student_2**.

b. Use appropriate attributes and methods to check that the **student_2** DataFrame contains the correct columns.

Documentation for ['select_dtypes'](#)

```

In [60]: student_2 = pd.read_table('student.txt', sep=';')
student_2.select_dtypes(include=['int64'])

```

```
Out[60]:
```

	age	Medu	Fedu	traveltime	studytime	failures	famrel	freetime
0	18	4	4	2	2	0	4	3
1	17	1	1	1	2	0	5	3
2	15	1	1	1	2	0	4	3
3	15	4	2	1	3	0	3	2
4	16	3	3	1	2	0	4	3
...
644	19	2	3	1	3	1	5	4
645	18	3	1	1	2	0	4	3
646	18	1	1	2	2	0	1	1
647	17	3	1	2	1	0	2	4
648	18	3	2	3	1	0	4	4

649 rows × 16 columns

PART FOUR

<div class="alert alert-block alert-info"

For Questions 17-20: We will be using the
'marketing_campaign.csv' dataset and the **customers DataFrame**
 </div>

Question 17:

- Read in the dataset **'marketing_campaign.csv'** and store the results in a DataFrame named **customers**.
- Use appropriate attributes and methods to inspect the **customers** DataFrame.

```
In [63]: customers = pd.read_csv('marketing_campaign.csv', sep=';')
print(customers.head())
print(customers.tail())
print(customers.info())
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome
0	5524	1957	Graduation	Single	58138.0	
1	2174	1954	Graduation	Single	46344.0	
2	4141	1965	Graduation	Together	71613.0	
3	6182	1984	Graduation	Together	26646.0	
4	5324	1981	PhD	Married	58293.0	

	Dt_Customer	Recency	MntWines	...	NumWebVisitsMonth	Accep
0	2012-09-04	58	635	...		7
1	2014-03-08	38	11	...		5
2	2013-08-21	26	426	...		4
3	2014-02-10	26	11	...		6
4	2014-01-19	94	173	...		5

	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Comp
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	Z_CostContact	Z_Revenue	Response
0	3	11	1
1	3	11	0
2	3	11	0
3	3	11	0
4	3	11	0

[5 rows x 29 columns]

	ID	Year_Birth	Education	Marital_Status	Income	Ki
2235	10870	1967	Graduation	Married	61223.0	
2236	4001	1946	PhD	Together	64014.0	
2237	7270	1981	Graduation	Divorced	56981.0	
2238	8235	1956	Master	Together	69245.0	
2239	9405	1954	PhD	Married	52869.0	

sMonth \	Teenhome	Dt_Customer	Recency	MntWines	...	NumWebVisit
2235	1	2013-06-13	46	709	...	
5						
2236	1	2014-06-10	56	406	...	
7						
2237	0	2014-01-25	91	908	...	
6						
2238	1	2014-01-24	8	428	...	
3						
2239	1	2012-10-15	40	84	...	
7						

AcceptedCmp2 \	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	A
2235	0	0	0	0	
0					
2236	0	0	0	1	
0					
2237	0	1	0	0	
0					
2238	0	0	0	0	
0					
2239	0	0	0	0	
0					

	Complain	Z_CostContact	Z_Revenue	Response
2235	0	3	11	0
2236	0	3	11	0
2237	0	3	11	0
2238	0	3	11	0
2239	0	3	11	1

[5 rows x 29 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2240 entries, 0 to 2239

Data columns (total 29 columns):

#	Column	Non-Null Count	Dtype
0	ID	2240 non-null	int64
1	Year_Birth	2240 non-null	int64
2	Education	2240 non-null	object
3	Marital_Status	2240 non-null	object
4	Income	2216 non-null	float64
5	Kidhome	2240 non-null	int64
6	Teenhome	2240 non-null	int64
7	Dt_Customer	2240 non-null	object
8	Recency	2240 non-null	int64
9	MntWines	2240 non-null	int64
10	MntFruits	2240 non-null	int64
11	MntMeatProducts	2240 non-null	int64
12	MntFishProducts	2240 non-null	int64
13	MntSweetProducts	2240 non-null	int64
14	MntGoldProds	2240 non-null	int64
15	NumDealsPurchases	2240 non-null	int64
16	NumWebPurchases	2240 non-null	int64
17	NumCatalogPurchases	2240 non-null	int64
18	NumStorePurchases	2240 non-null	int64
19	NumWebVisitsMonth	2240 non-null	int64
20	AcceptedCmp3	2240 non-null	int64

```
21 AcceptedCmp4      2240 non-null  int64
22 AcceptedCmp5      2240 non-null  int64
23 AcceptedCmp1      2240 non-null  int64
24 AcceptedCmp2      2240 non-null  int64
25 Complain          2240 non-null  int64
26 Z_CostContact      2240 non-null  int64
27 Z_Revenue          2240 non-null  int64
28 Response           2240 non-null  int64
```

```
dtypes: float64(1), int64(25), object(3)
```

```
memory usage: 507.6+ KB
```

```
None
```

Documentation for [describe\(\)](#)

Question 18:


Use the appropriate method to generate summary statistics for the numeric columns of the **customers** DataFrame.

```
In [64]: customers.describe()
```

```
Out[64]:
```

	ID	Year_Birth	Income	Kidhome	Teenhome
count	2240.000000	2240.000000	2216.000000	2240.000000	2240.000000
mean	5592.159821	1968.805804	52247.251354	0.444196	0.506250
std	3246.662198	11.984069	25173.076661	0.538398	0.544538
min	0.000000	1893.000000	1730.000000	0.000000	0.000000
25%	2828.250000	1959.000000	35303.000000	0.000000	0.000000
50%	5458.500000	1970.000000	51381.500000	0.000000	0.000000
75%	8427.750000	1977.000000	68522.000000	1.000000	1.000000
max	11191.000000	1996.000000	666666.000000	2.000000	2.000000

8 rows × 6 columns



Question 19:

Use the appropriate method to generate summary statistics for the non-numeric columns of the **customers** DataFrame.

```
In [70]: customers.describe(include=['object'])
```

Out[70]:

	Education	Marital_Status	Dt_Customer
count	2240	2240	2240
unique	5	8	663
top	Graduation	Married	2012-08-31
freq	1127	864	12

Question 20:

Use the appropriate method to generate summary statistics for the all columns of the **customers** DataFrame.

In [72]: `customers.describe(include=['object', np.number])`

Out[72]:

	ID	Year_Birth	Education	Marital_Status	In
count	2240.000000	2240.000000	2240	2240	2216.0
unique	NaN	NaN	5	8	
top	NaN	NaN	Graduation	Married	
freq	NaN	NaN	1127	864	
mean	5592.159821	1968.805804	NaN	NaN	52247.2
std	3246.662198	11.984069	NaN	NaN	25173.0
min	0.000000	1893.000000	NaN	NaN	1730.0
25%	2828.250000	1959.000000	NaN	NaN	35303.0
50%	5458.500000	1970.000000	NaN	NaN	51381.5
75%	8427.750000	1977.000000	NaN	NaN	68522.0
max	11191.000000	1996.000000	NaN	NaN	666666.0

11 rows × 29 columns

PART FIVE

<div class="alert alert-block alert-info">

For Questions 17-20: We will be using the '**cereals.txt**' **dataset** and the **cereals DataFrame** </div>

Question 21:

a. Read in the dataset '**cereals.txt**' and store the results in a DataFrame named **cereals**.

See the link below for access to the dataset.

[cereals.txt](#)

b. Use appropriate attributes and methods to inspect the **cereals** DataFrame.

```
In [75]: cereals = pd.read_table('cereals.txt', sep=',')
cereals
```

```
Out[75]:
```

	Name	Manufacturer	Type	Calories	Fiber	Sugars
0	100% Bran	Nabisco	Cold	70	10.0	6
1	100% Natural Bran	Quaker Oats	Cold	120	2.0	8
2	All-Bran	Kellogg's	Cold	70	9.0	5
3	All-Bran with Extra Fiber	Kellogg's	Cold	50	14.0	0
4	Almond Delight	Ralston Purina	Cold	110	1.0	8
...
72	Triples	General Mills	Cold	110	0.0	3
73	Trix	General Mills	Cold	110	0.0	12
74	Wheat Chex	Ralston Purina	Cold	100	3.0	3
75	Wheaties	General Mills	Cold	100	3.0	3
76	Wheaties Honey Gold	General Mills	Cold	110	1.0	8

77 rows × 6 columns

Question 22:

a. Use the **drop()** method to temporarily drop the "Type" column of the **cereals** DataFrame.

 b. Run the **cereals.head()** command. What do you notice?

Documentation for 'drop'

```
In [82]: cereals.drop(["Type"],axis=1)
print("The type is still there because we did not use th
cereals.head()
```

The type is still there because we did not use the inplace argument.

```
Out[82]:
```

	Name	Manufacturer	Type	Calories	Fiber	Sugars
0	100% Bran	Nabisco	Cold	70	10.0	6
1	100% Natural Bran	Quaker Oats	Cold	120	2.0	8
2	All-Bran	Kellogg's	Cold	70	9.0	5
3	All-Bran with Extra Fiber	Kellogg's	Cold	50	14.0	0
4	Almond Delight	Ralston Purina	Cold	110	1.0	8

Question 23:

a. Use the **drop()** method to temporarily drop the third row of the **cereals**

 b. Run the **cereals.head()** command. What do you notice?

Documentation for 'drop'

```
In [84]: cereals.drop(2,axis=0)
print("The third row is still there because we did not
cereals.head()
```

The type is still there because we did not use the inplace argument.

Out[84]:

	Name	Manufacturer	Type	Calories	Fiber	Sugars
0	100% Bran	Nabisco	Cold	70	10.0	6
1	100% Natural Bran	Quaker Oats	Cold	120	2.0	8
2	All-Bran	Kellogg's	Cold	70	9.0	5
3	All-Bran with Extra Fiber	Kellogg's	Cold	50	14.0	0
4	Almond Delight	Ralston Purina	Cold	110	1.0	8

Question 24:

Calculate the mean of each numeric column.

Documentation for ['mean'](#)

In [87]: `cereals.mean(numeric_only=True)`

Out[87]:

```
Calories    106.883117
Fiber        2.151948
Sugars       6.922078
dtype: float64
```

Question 25:

Calculate the mean of each row.

Documentation for ['mean'](#)

In [90]: `cereals.mean(axis=1, numeric_only=True)`

Out[90]:

```
0      28.666667
1      43.333333
2      28.000000
3      21.333333
4      39.666667
...
72      37.666667
73      40.666667
74      35.333333
75      35.333333
76      39.666667
Length: 77, dtype: float64
```

Question 26:

Calculate the mean of each numeric column using the alias "index" for axis=0.

```
In [91]: cereals.mean(axis=0, numeric_only=True)
```

```
Out[91]: Calories    106.883117  
Fiber           2.151948  
Sugars          6.922078  
dtype: float64
```

Question 27:

Calculate the mean of each row using the alias "columns" for axis=1.

```
In [92]: cereals.mean(axis=1, numeric_only=True)
```

```
Out[92]: 0      28.666667  
1      43.333333  
2      28.000000  
3      21.333333  
4      39.666667  
...  
72     37.666667  
73     40.666667  
74     35.333333  
75     35.333333  
76     39.666667  
Length: 77, dtype: float64
```